

Jesus Hernandez: Worked on Part 2 of the project and the regression part of Part 3
Isaac Jee: Worked on the classification part of Part 3
Jake Takahashi: Worked on part 1 and 2 and did some minor debugging on Part 3
Aaron Castiglione: Worked on part 1 and part 4.

Sources:

<https://docs.h2o.ai/h2o/latest-stable/h2o-docs/data-science/gbm.html>

https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html

<https://www.ml-science.com/masking>

<https://stackoverflow.com/questions/47216388/efficiently-creating-masks-numpy-python>

<https://developers.google.com/machine-learning/crash-course/overfitting/regularization>

<https://xgboost.readthedocs.io/en/stable/tutorials/model.html>

<https://www.geeksforgeeks.org/machine-learning/xgboost/>

<https://medium.com/data-science/pca-102-should-you-use-pca-how-many-components-to-use-how-to-interpret-them-da0c8e3b11f0>

<https://www.geeksforgeeks.org/machine-learning/stratified-k-fold-cross-validation/>

<https://www.geeksforgeeks.org/machine-learning/cross-validation-machine-learning/>